



UTM

UNIVERSITI TEKNOLOGI MALAYSIA

FACULTY OF COMPUTING
SEMESTER II, SESSION 2024/2025

BACHELOR OF COMPUTER SCIENCE (BIOINFORMATICS)

SECB3203 PROGRAMMING FOR BIOINFORMATICS - SECTION 01

**PROJECT
ALZHEIMER'S DISEASE PREDICTION**

GROUP MEMBERS	MATRIC NO
TAN ZHAO HONG	A23CS0188
CHIN PEI WEN	A23CS0065
KOO XUAN	A23CS0300

LECTURER'S NAME : DR. SEAH CHOON SEN

SUBMISSION DATE : 03 DECEMBER 2025

Table of Contents

2.0 Data Collection and Pre-processing.....	3
2.1 Importing Dataset.....	4
- Python packages for Data Science.....	4
- Importing Dataset File.....	5
- Dataset Information.....	6
- Feature Removal.....	7
2.2 Data wrangling.....	8
- Target Variable Distribution.....	8
- Identifying and handling missing values.....	9
- Data Formatting.....	10
- Data Normalization.....	11
- Binning.....	13
- Indicator variables.....	15
- Train–Test Split and Z-Score Normalization.....	16

2.0 Data Collection and Pre-processing

Our dataset was fetched from Kaggle:

<https://www.kaggle.com/datasets/rabieelkharoua/alzheimers-disease-dataset/data>

The dataset is in .xlsx, thus we convert it to csv file so that it can be easily read by Python.

2.1 Importing Dataset

- Python packages for Data Science

```
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
9 from sklearn.preprocessing import StandardScaler, LabelEncoder
10 from sklearn.metrics import (accuracy_score, precision_score, recall_score,
11 | | | | | | | | | | | | f1_score, confusion_matrix, classification_report,
12 | | | | | | | | | | | | roc_auc_score, roc_curve)
13 # from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.svm import SVC
16 from sklearn.tree import DecisionTreeClassifier
17 from sklearn.neighbors import KNeighborsClassifier
18 import warnings
19 warnings.filterwarnings('ignore')
```

- pandas: Data manipulation and analysis
- numpy: Numerical operations and array handling
- matplotlib: Data visualization
- seaborn: Statistical data visualization
- scikit-learn: Machine learning algorithms and preprocessing

- Importing Dataset File

```
37  # =====
38  # 1. DATA LOADING
39  # =====
40
41 ✓ def load_data(filepath):
42     """Load the Alzheimer's disease dataset"""
43     df = pd.read_csv(filepath)
44     print(f"Dataset shape: {df.shape}")
45     print(df)
46
| 367 ✓ if __name__ == "__main__":
368     FILEPATH = "alzheimers_disease_data.csv"
369     TARGET_COLUMN = 'Diagnosis'
=====
ALZHEIMER'S DISEASE PREDICTION - ML PIPELINE
=====
[1] Loading Data...
Dataset shape: (2149, 35)
   PatientID  Age  Gender  Ethnicity  EducationLevel      BMI ... Disorientation  PersonalityChanges  DifficultyCompletingTasks  Forgetfulness  Diagnosis  DoctorInCharge
0       4751    73      0          0            2  22.927749 ...           0             0               1              0              0                0        XXXConfid
1       4752    89      0          0            0  26.827681 ...           0             0               0              1              0                0        XXXConfid
2       4753    73      0          3            1  17.795882 ...           1             0               1              0              0                0        XXXConfid
3       4754    74      1          0            1  33.800817 ...           0             0               0              0              0                0        XXXConfid
4       4755    89      0          0            0  20.716974 ...           0             0               1              1              0                0        XXXConfid
...
...   ...   ...
2144    6895    61      0          0            1  39.121757 ...           0             0               0              0              1                1        XXXConfid
2145    6896    75      0          0            2  17.857903 ...           0             0               0              0              1                1        XXXConfid
2146    6897    77      0          0            1  15.476479 ...           0             0               0              0              1                1        XXXConfid
2147    6898    78      1          3            1  15.299911 ...           0             0               0              0              1                1        XXXConfid
2148    6899    72      0          0            2  33.289738 ...           1             0               0              0              1                0        XXXConfid
[2149 rows x 35 columns]
```

These screenshots show the initial dataset loading process.

- Dataset Information

```
46 |     print(f"\nDataset info:")
47 |     print(df.info())
```

```
Dataset info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2149 entries, 0 to 2148
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   PatientID        2149 non-null    int64  
 1   Age              2149 non-null    int64  
 2   Gender            2149 non-null    int64  
 3   Ethnicity         2149 non-null    int64  
 4   EducationLevel   2149 non-null    int64  
 5   BMT              2149 non-null    float64 
 6   Smoking           2149 non-null    int64  
 7   AlcoholConsumption 2149 non-null    float64 
 8   PhysicalActivity  2149 non-null    float64 
 9   DietQuality       2149 non-null    float64 
 10  SleepQuality      2149 non-null    float64 
 11  FamilyHistoryAlzheimers 2149 non-null    int64  
 12  CardiovascularDisease 2149 non-null    int64  
 13  Diabetes          2149 non-null    int64  
 14  Depression         2149 non-null    int64  
 15  HeadInjury         2149 non-null    int64  
 16  Hypertension        2149 non-null    int64  
 17  SystolicBP         2149 non-null    int64  
 18  DiastolicBP        2149 non-null    int64  
 19  CholesterolTotal   2149 non-null    float64 
 20  CholesterolLDL    2149 non-null    float64 

 21  CholesterolHDL    2149 non-null    float64 
 22  CholesterolTriglycerides 2149 non-null    float64 
 23  MMSE              2149 non-null    float64 
 24  FunctionalAssessment 2149 non-null    float64 
 25  MemoryComplaints  2149 non-null    int64  
 26  BehavioralProblems 2149 non-null    int64  
 27  ADL               2149 non-null    float64 
 28  Confusion          2149 non-null    int64  
 29  Disorientation      2149 non-null    int64  
 30  PersonalityChanges 2149 non-null    int64  
 31  DifficultyCompletingTasks 2149 non-null    int64  
 32  Forgetfulness       2149 non-null    int64  
 33  Diagnosis          2149 non-null    int64  
 34  DoctorInCharge     2149 non-null    object  
dtypes: float64(12), int64(22), object(1)
memory usage: 587.7+ KB
None
```

These screenshots show the dataset information details.

- Feature Removal


```
25 # =====
26 # FEATURES TO REMOVE (Non-medical and weak evidence features)
27 # =====
28 FEATURES_TO_REMOVE = [
29     'PatientID',           # Administrative identifier
30     'DoctorInCharge',      # Administrative data
31     'Gender',              # Weak predictor
32     'Ethnicity',           # May introduce bias, weak medical relevance
33     'AlcoholConsumption',  # Inconsistent evidence
34     'SleepQuality'         # Weak/inconsistent evidence
35 ]
36
37
38 # Remove non-medical and weak evidence features
39 print(f"\n{'='*80}")
40 print("REMOVING NON-MEDICAL AND WEAK EVIDENCE FEATURES")
41 print(f"{'='*80}")
42 features_removed = [col for col in FEATURES_TO_REMOVE if col in df.columns]
43 if features_removed:
44     print(f"\nRemoving features: {features_removed}")
45     df = df.drop(columns=features_removed)
46     print(f"New dataset shape: {df.shape}")
47 else:
48     print("\nNo features to remove (features not found in dataset)")
49
50 return df
51
52 =====
53 REMOVING NON-MEDICAL AND WEAK EVIDENCE FEATURES
54 =====
55
56 Removing features: ['PatientID', 'DoctorInCharge', 'Gender', 'Ethnicity', 'AlcoholConsumption', 'SleepQuality']
57 New dataset shape: (2149, 29)
```

These screenshots show the removal of non-medical and weak-evidence features such as patient identifiers and administrative attributes. These features were excluded to reduce noise, prevent bias, and ensure that the model focuses only on medically relevant variables that contribute meaningfully to Alzheimer's disease prediction.

2.2 Data wrangling

- Target Variable Distribution

```
110     def preprocess_data(df, target_column='Diagnosis'):
111         """Preprocess the data for machine learning"""
112
113         # Separate features and target
114         X = df.drop(columns=[target_column])
115         y = df[target_column]
116
117         # Encode target variable if it's categorical
118         if y.dtype == 'object':
119             print(f"\nEncoding target variable")
120             le_target = LabelEncoder()
121             y = le_target.fit_transform(y)
```



This figure illustrates the distribution of the target variable (Diagnosis) using both bar and pie charts. Visualizing the class distribution helps identify potential class imbalance and confirms that the problem is a binary classification task, which is suitable for the selected machine learning models.

- Identifying and handling missing values


```

123 # VISUALIZE MISSING VALUES
124 print("\n--- Identifying and Handling Missing Values ---")
125 print("\n1. Missing Values Count:")
126 print(x.isnull().sum())
127
128 print("\n2. Visualizing Missing Values:")
129 plt.figure(figsize=(12, 6))
130 sns.heatmap(x.isnull(), cbar=True, cmap='viridis', yticklabels=False)
131 plt.title('Missing Values Heatmap', fontsize=14, fontweight='bold')
132 plt.xlabel('Columns')
133 plt.tight_layout()
134 plt.savefig('missing_values_heatmap.png', dpi=300)
135 plt.show()
136
137 # Handle missing values
138 if x.isnull().sum().sum() > 0:
139     print("\n3. Handling missing values...")
140     x = x.fillna(x.median())
141     print("Missing values after handling:")
142     print(x.isnull().sum())
143 else:
144     print("\n3. No missing values detected!")

```

2. Visualizing Missing Values:

3. No missing values detected!



These screenshots display the detection of missing values across all input features. Median imputation was applied to numerical attributes to handle missing data, as it is robust to outliers and preserves the overall distribution of the dataset. This step ensures data completeness before further processing.

- Data Formatting

```
146     # DATA FORMATTING
147     print("\n--- Data Formatting ---")
148     print("\n1. Current Data Types:")
149     print(x.dtypes)
150
151     print("\n2. Numerical Columns:")
152     numerical_cols = x.select_dtypes(include=[np.number]).columns
153     print(list(numerical_cols))
154
155     print("\n3. Categorical Columns:")
156     categorical_cols_orig = x.select_dtypes(include=['object']).columns
157     print(list(categorical_cols_orig))
```

```
--- Data Formatting ---
1. Current Data Types:
Age                      int64
EducationLevel            int64
BMT                      float64
Smoking                  int64
PhysicalActivity          float64
DietQuality              float64
FamilyHistoryAlzheimers   int64
CardiovascularDisease    int64
Diabetes                 int64
Depression               int64
HeadInjury                int64
Hypertension              int64
SystolicBP                int64
DiastolicBP              int64
CholesterolTotal          float64
CholesterolLDL           float64
CholesterolHDL           float64
CholesterolTriglycerides float64
MMSE                     float64
FunctionalAssessment      float64
MemoryComplaints          int64
BehavioralProblems        int64
ADL                      float64
Confusion                 int64
Disorientation             int64
PersonalityChanges         int64
DifficultyCompletingTasks int64
Forgetfulness              int64
dtype: object

2. Numerical Columns:
['Age', 'EducationLevel', 'BMT', 'Smoking', 'PhysicalActivity', 'DietQuality', 'FamilyHistoryAlzheimers', 'cardiovascularDisease', 'Diabetes', 'Depression', 'HeadInjury', 'Hypertension', 'SystolicBP', 'DiastolicBP', 'CholesterolTotal', 'CholesterolLDL', 'CholesterolHDL', 'CholesterolTriglycerides', 'MMSE', 'FunctionalAssessment', 'MemoryComplaints', 'BehavioralProblems', 'ADL', 'Confusion', 'Disorientation', 'PersonalityChanges', 'DifficultyCompletingTasks', 'Forgetfulness']

3. Categorical columns:
[]
```

This step verifies and formats the dataset by inspecting data types and separating numerical and categorical features. Ensuring correct data types is essential before applying normalization, binning, and encoding, as machine learning algorithms require numerical and consistently formatted inputs.

- Data Normalization

```
159     # DATA NORMALIZATION
160     print("\n--- Data Normalization ---")
161
162     # Import MinMaxScaler
163     from sklearn.preprocessing import MinMaxScaler
164
165     print("\n1. Min-Max Scaling (Normalization)")
166     print("  Formula:  $x_{scaled} = (x - x_{min}) / (x_{max} - x_{min})$ ")
167     print("  Result: Values scaled to range [0, 1]")
168
169     if 'Age' in X.columns:
170         # Create a copy for demonstration
171         min_max_scaler = MinMaxScaler()
172         age_minmax = min_max_scaler.fit_transform(X[['Age']])
173
174         print("\n  Example: Age Column")
175         print(f"  Original range: [{X['Age'].min():.2f}, {age_minmax.min():.2f}]")
176         print(f"  Scaled range: [{age_minmax.min():.2f}, {age_minmax.max():.2f}]")
177
178
179     # Show comparison
180     comparison_df = pd.DataFrame({
181         'Original_Age': X['Age'].head(10),
182         'MinMax_Scaled': age_minmax[:10].flatten()
183     })
184     print("\n  Sample comparison:")
185     print(comparison_df)
186
187     # Visualize
188     fig, axes = plt.subplots(1, 2, figsize=(14, 5))
189     axes[0].hist(X['Age'], bins=20, color='skyblue', edgecolor='black')
190     axes[0].set_title('Original Age Distribution', fontsize=12, fontweight='bold')
191     axes[0].set_xlabel('Age')
192     axes[0].set_ylabel('Frequency')
193
194     axes[1].hist(age_minmax, bins=20, color='salmon', edgecolor='black')
195     axes[1].set_title('Min-Max Scaled Age [0,1]', fontsize=12, fontweight='bold')
196     axes[1].set_xlabel('Scaled Age')
197     axes[1].set_ylabel('Frequency')
198
199     plt.tight_layout()
200     plt.savefig('minmax_scaling.png', dpi=300)
201     plt.show()
202
203     print("\n2. Z-Score Normalization (Standardization)")
204     print("  Formula:  $z = (x - \mu) / \sigma$ ")
205     print("  Note: This will be applied later using StandardScaler")
```

--- Data Normalization ---

1. Min-Max Scaling (Normalization)

Formula: $x_{scaled} = (x - x_{min}) / (x_{max} - x_{min})$

Result: Values scaled to range [0, 1]

Example: Age column

Original range: [60.00, 90.00]

Scaled range: [0.00, 1.00]

Sample comparison:


Original_Age MinMax_Scaled

0	73	0.433333
1	89	0.966667
2	73	0.433333
3	74	0.466667
4	89	0.966667
5	86	0.866667
6	68	0.266667
7	75	0.500000
8	72	0.400000
9	87	0.900000

2. Z-Score Normalization (Standardization)

Formula: $z = (x - \mu) / \sigma$

Note: This will be applied later using StandardScaler



This output demonstrates Min-Max normalization applied to the Age feature, scaling values into the range [0,1]. The comparison between original and scaled values highlights how normalization prevents features with larger numerical ranges from dominating the learning process of distance-based and gradient-based models.

- Binning

```
206 # BINNING
207 print("\n--- Binning ---")
208
209 if 'Age' in X.columns:
210     print("\n1. Age Binning:")
211     X['Age_binned'] = pd.cut(
212         X['Age'],
213         bins=[0, 60, 75, 100],
214         labels=['Young', 'Middle', 'Old']
215     )
216
217     # Show binning results
218     print("\n    Sample of Age Binning:")
219     print(X[['Age', 'Age_binned']].head(10))
220
221     print("\n    Age Group Distribution:")
222     print(X['Age_binned'].value_counts())
223
224     # Visualize
225     plt.figure(figsize=(10, 6))
226     X['Age_binned'].value_counts().sort_index().plot(kind='bar', color='steelblue', edgecolor='black')
227     plt.title('Age Group Distribution After Binning', fontsize=14, fontweight='bold')
228     plt.xlabel('Age Group')
229     plt.ylabel('Count')
230     plt.xticks(rotation=0)
231     plt.tight_layout()
232     plt.savefig('age_binning.png', dpi=300)
233     plt.show()
```

--- Binning ---

1. Age Binning:


Sample of Age Binning:

	Age	Age_binned
0	73	Middle
1	89	Old
2	73	Middle
3	74	Middle
4	89	Old
5	86	Old
6	68	Middle
7	75	Middle
8	72	Middle
9	87	Old

Age Group Distribution:

Age_binned	count
Middle	1048
Old	1027
Young	74

Name: count, dtype: int64



In this step, the continuous Age feature was transformed into categorical age groups using binning. This approach helps capture non-linear patterns related to age and Alzheimer's risk, improves interpretability, and allows the model to learn differences between broader age groups rather than relying only on exact numerical values.

- Indicator variables

```
235 # INDICATOR VARIABLES
236 print("\n--- Indicator Variables (One-Hot Encoding) ---")
237
238 categorical_cols = X.select_dtypes(include=['object', 'category']).columns
239
240 if len(categorical_cols) > 0:
241     print(f"\n1. Categorical columns to encode: {list(categorical_cols)}")
242
243     # Show before encoding
244     print("\n2. Before encoding (first 5 rows):")
245     print(X[categorical_cols].head())
246
247     # Apply one-hot encoding
248     X = pd.get_dummies(X, columns=categorical_cols, drop_first=True)
249
250     print(f"\n3. After one-hot encoding:")
251     print(f"    Original features: {len(categorical_cols)}")
252     print(f"    New indicator columns created: {len([col for col in X.columns if any(cat in col for cat in categorical_cols)])}")
253     print(f"    Total columns now: {X.shape[1]}")
```

--- Indicator Variables (One-Hot Encoding) ---

1. Categorical columns to encode: ['Age_binned']

2. Before encoding (first 5 rows):

Age_binned

0	Middle
1	Old
2	Middle
3	Middle
4	Old

3. After one-hot encoding:

Original features: 1

New indicator columns created: 2

Total columns now: 30

These screenshots show the conversion of categorical features created through binning into numerical indicator variables using one-hot encoding. This transformation is necessary because machine learning models require numerical inputs and cannot directly process categorical text labels.

- Train–Test Split and Z-Score Normalization

```
254
255     # SPLIT AND SCALE
256     # Split the data
257     X_train, X_test, y_train, y_test = train_test_split(
258         X, y, test_size=0.2, random_state=42, stratify=y
259     )
260
261     # Scale the features (Z-score normalization)
262     print("\n--- Final Scaling (Z-Score Normalization) ---")
263     scaler = StandardScaler()
264     X_train_scaled = scaler.fit_transform(X_train)
265     X_test_scaled = scaler.transform(X_test)
266
267     print(f"\nTraining set size: {X_train_scaled.shape}")
268     print(f"Test set size: {X_test_scaled.shape}")
269
270     print("\n" + "="*80)
271     print("DATA PREPROCESSING COMPLETE")
272     print("="*80)
273
274     return X_train_scaled, X_test_scaled, y_train, y_test, X.columns
275
```

```
--- Final Scaling (Z-Score Normalization) ---
```

```
Training set size: (1719, 30)
Test set size: (430, 30)
```

```
=====
DATA PREPROCESSING COMPLETE
=====
```

The dataset was split into training and testing sets using stratified sampling to preserve class distribution. Z-score normalization was then applied using StandardScaler to standardize features, ensuring that all variables contribute equally during model training.